

NYT Prediction Report

Exploring Data

There are 8402 observations in total, where 6532 training observations and 1870 observations.

```
popularDensity <- table(newsData$Popular)
posPopular <- round(popularDensity[2]/(popularDensity[1] + popularDensity[2]) * 100, 2)
```

Only 16.73% of all New York Times blog articles have more than 25 comments. That means, a baseline model for predicting unpopular would be around 83.27%.

The independent variables consist of 8 pieces of article data available at the time of publication, and a unique identifier:

- **NewsDesk**, the New York Times desk that produced the story (Business, Culture, Foreign, etc.)
- **SectionName**, the section the article appeared in (Opinion, Arts, Technology, etc.)
- **SubsectionName**, the subsection the article appeared in (Education, Small Business, Room for Debate, etc.)
- **Headline**, the title of the article
- **Snippet**, a small portion of the article text
- **Abstract**, a summary of the blog article, written by the New York Times
- **WordCount**, the number of words in the article
- **PubDate**, the publication date, in the format “Year-Month-Day Hour:Minute:Second”
- **UniqueID**, a unique identifier for each article

Cleaning Data

Text of Articles

Let's take a look at **Headline**, we can observe many very common combination of words like `new york times`, `pictures of the day` etc. If we google `pictures of the day new york times`, it is easy to know `pictures of the day` is a daily article from `Lens` category.

```
newsData$Headline
```


To avoid overcounting of words like `day`, a replacement of some proper nouns to single word is necessary.

```
originalText = c("new york times", "new york city", "new york", "silicon valley",
                 "times insider", "fashion week", "white house",
                 "international herald tribune archive",
                 "president obama", "hong kong", "big data", "golden globe",
                 "word of the day", "time square", "pictures of the day",
                 "photos of the day", "daily clip report", "daily report",
                 "6 q's about the news", "test yourself")

replacementText = c("NYT", "NYC", "NewYork", "SiliconValley", "TimesInsider",
                   "FashionWeek", "WhiteHouse", "IHT", "Obama", "HongKong",
                   "BigData", "GoldenGlobe", "WordofDay", "TimeSquare", "PicOfDay",
                   "PicOfDay", "DailyClipReport", "DailyReport", "6q", "TestYourself")
```

Word Count of Articles

The following plot shows article's popularity distribution based on logarithmic word count. Beside training data, testing data's distribution is also plotted by gray color which is bimodal distribution.



If we conduct a two-sided t-test on the mean and a two-sided F-test on the variance:

```
PopularNewsTrain = subset(newsTrain, newsTrain$Popular==1)
UnpopularNewsTrain = subset(newsTrain, newsTrain$Popular==0)

t.test(PopularNewsTrain$LogWordCount, UnpopularNewsTrain$LogWordCount)


##
## Welch Two Sample t-test
##
## data: PopularNewsTrain$LogWordCount and UnpopularNewsTrain$LogWordCount
## t = 28.2691, df = 2310.719, p-value < 2.2e-16
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##  0.8018917 0.9214367
## sample estimates:
## mean of x mean of y
##  6.455548 5.593884
```

```
var.test(PopularNewsTrain$LogWordCount, UnpopularNewsTrain$LogWordCount)
```

```
##
## F test to compare two variances
##
```

```
## data: PopularNewsTrain$LogWordCount and UnpopularNewsTrain$LogWordCount
## F = 0.4108, num df = 1092, denom df = 5438, p-value < 2.2e-16
## alternative hypothesis: true ratio of variances is not equal to 1
## 95 percent confidence interval:
## 0.3752262 0.4509859
## sample estimates:
## ratio of variances
## 0.4107796
```

This shows us there is a statistically significant difference between popular and unpopular articles based on the word counts. At the same time, popular article seems having shorter **Headline**.



```
t.test(PopularNewsTrain$HeadlineCharCount, UnpopularNewsTrain$HeadlineCharCount)
```

```
##
## Welch Two Sample t-test
##
## data: PopularNewsTrain$HeadlineCharCount and UnpopularNewsTrain$HeadlineCharCount
## t = -10.4608, df = 1483.142, p-value < 2.2e-16
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -7.191321 -4.920218
## sample estimates:
## mean of x mean of y
## 41.16651 47.22228
```

Publishing Hour and Day

It is unlikely that many article receiving 25 more comments in the middle of night. Hence, at certain times during the day, we expect the probability that a random article becomes popular to be larger. Similarly, the day of the week may have an impact, people may have much more time to read articles than a working day.

```

## date feature
newsData$PubDate = strptime(newsData$PubDate, "%Y-%m-%d %H:%M:%S")
newsData$PubDay = as.Date(newsData$PubDate)
## it is expected that different behaviours at different times of the day.publication
newsData$DayOfWeek = newsData$PubDate$yday
newsData$Hour = newsData$PubDate$hour
newsData$DayOfWeek = as.factor(weekdays(newsData$PubDate))
newsData$DayOfWeek = factor(newsData$DayOfWeek, levels=c("Monday", "Tuesday", "Wednesday", "Thursday",


## especially on holidays, people may have much more time to read and comment on blog articles
Holidays = c(as.POSIXlt("2014-09-01 00:00", format="%Y-%m-%d %H:%M"),
  as.POSIXlt("2014-10-13 00:00", format="%Y-%m-%d %H:%M"),
  as.POSIXlt("2014-10-31 00:00", format="%Y-%m-%d %H:%M"),
  as.POSIXlt("2014-11-11 00:00", format="%Y-%m-%d %H:%M"),
  as.POSIXlt("2014-11-27 00:00", format="%Y-%m-%d %H:%M"),
  as.POSIXlt("2014-12-24 00:00", format="%Y-%m-%d %H:%M"),
  as.POSIXlt("2014-12-25 00:00", format="%Y-%m-%d %H:%M"),
  as.POSIXlt("2014-12-31 00:00", format="%Y-%m-%d %H:%M"))

newsData$Holiday = as.factor(ifelse(newsData$PubDate$yday %in% Holidays$yday, 1, 0))

```


##	Unpopular	Popular	PopularDensity	Hour
## 4	169	2	0.01169591	4
## 5	240	11	0.04382470	5
## 3	57	3	0.05000000	3
## 7	369	25	0.06345178	7
## 1	25	3	0.10714286	1
## 8	299	39	0.11538462	8
## 2	21	3	0.12500000	2
## 13	390	57	0.12751678	13
## 17	348	52	0.13000000	17
## 6	190	30	0.13636364	6
## 9	267	49	0.15506329	9
## 16	382	74	0.16228070	16
## 14	375	77	0.17035398	14
## 0	103	22	0.17600000	0
## 10	315	69	0.17968750	10
## 12	428	94	0.18007663	12
## 18	269	60	0.18237082	18
## 11	414	93	0.18343195	11
## 15	345	93	0.21232877	15
## 19	140	53	0.27461140	19
## 20	114	44	0.27848101	20
## 21	76	34	0.30909091	21
## 23	46	21	0.31343284	23
## 22	57	85	0.59859155	22

It seems that publishing blog posts around 10 pm are more easier getting popular according to PopularDensity. But, if we compare number of popular articles of 24 hours, It is clear to see that around 12 pm and 3 pm, even more blog posts receiving 25 more comments than 10 pm.



##	Unpopular	Popular	PopularDensity	Day
## Friday	1003	161	0.1383162	5
## Tuesday	1016	174	0.1462185	2
## Thursday	1033	195	0.1587948	4
## Wednesday	1027	197	0.1609477	3
## Monday	1010	214	0.1748366	1
## Saturday	138	52	0.2736842	6
## Sunday	212	100	0.3205128	7

Similarly, day of week shows same trends as hourly results. Also, much more articles are published on weekday than weekends.



Category of Articles

There are three variables categorizes blog posts, NewsDesk SectionName and SubsectionName.

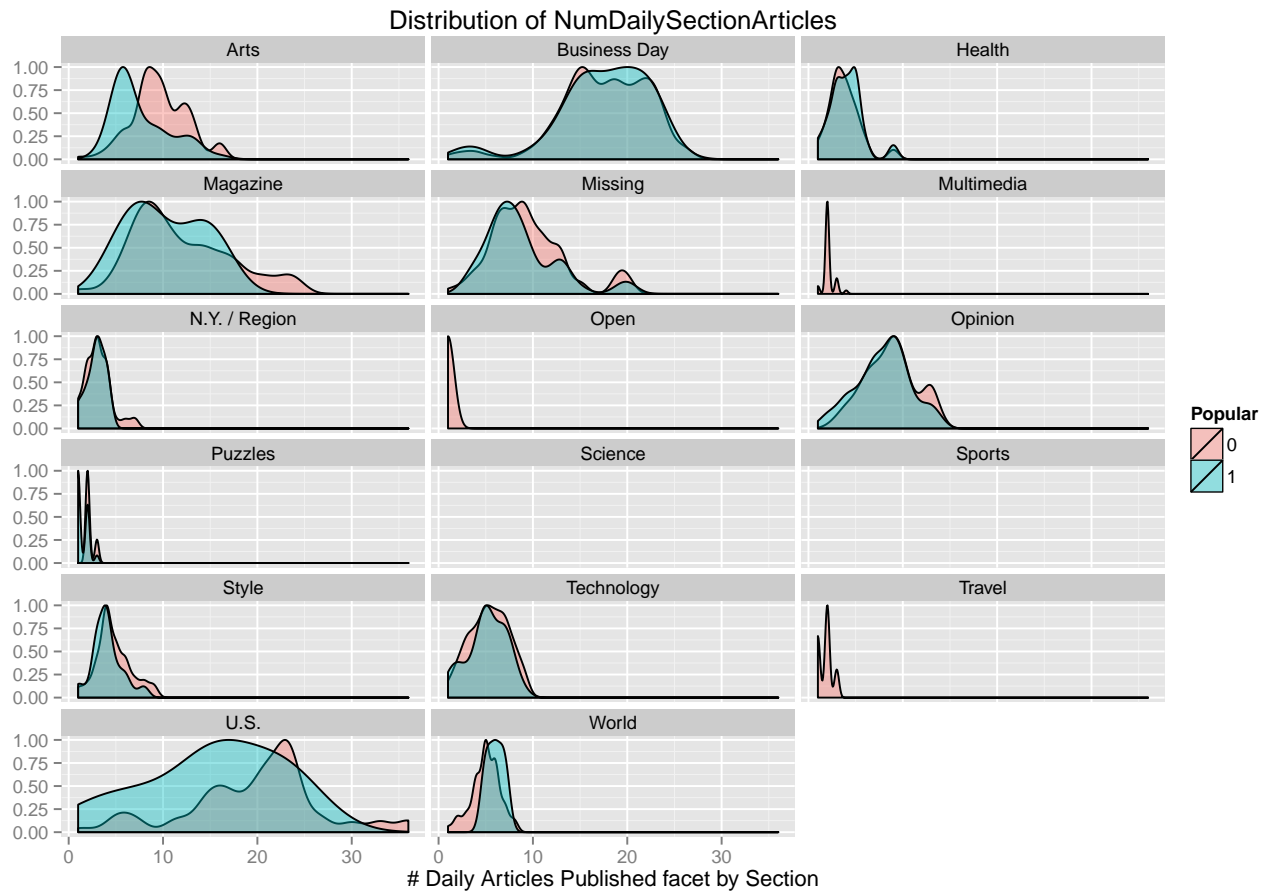
```
## missing categories
misCategory = subset(newsData, newsData$NewsDesk==" " | newsData$SectionName==" " | newsData$SubsectionName==" ")
dim(misCategory)[1]
```

```
## [1] 6721
```


```
misCategory = subset(newsData, newsData$NewsDesk==" " & newsData$SectionName==" " & newsData$SubsectionName==" ")
dim(misCategory)[1]
```

```
## [1] 1626
```

6721 articles have at one category variable missing and **1626** articles have no categories at all. After filling blank categories based on existing category variables, let's try to see the facet distribution of blog posts.



Although it is hard to see what's going on, a clear difference between popular and unpopular articles is in the section Magazine, where around 15 articles posted per day is more indicative of popular articles than unpopular ones. Beyond 20 posts per day the roles are clearly reversed. Hourly distribution of the following plot also shows no clear indication of popularity.



Contents Features of Articles

Until now, we preprocessed date features, word counts and categories of article. Almost all features from original data frame but the contents of blog post.

```

stopWords = c(stopwords("SMART"))
CorpusText = Corpus(VectorSource(newsData$Headline))
CorpusText = tm_map(CorpusText, tolower)
CorpusText = tm_map(CorpusText, PlainTextDocument)
CorpusText = tm_map(CorpusText, removePunctuation)
CorpusText = tm_map(CorpusText, removeWords, stopWords)
CorpusText = tm_map(CorpusText, stemDocument, language="english")

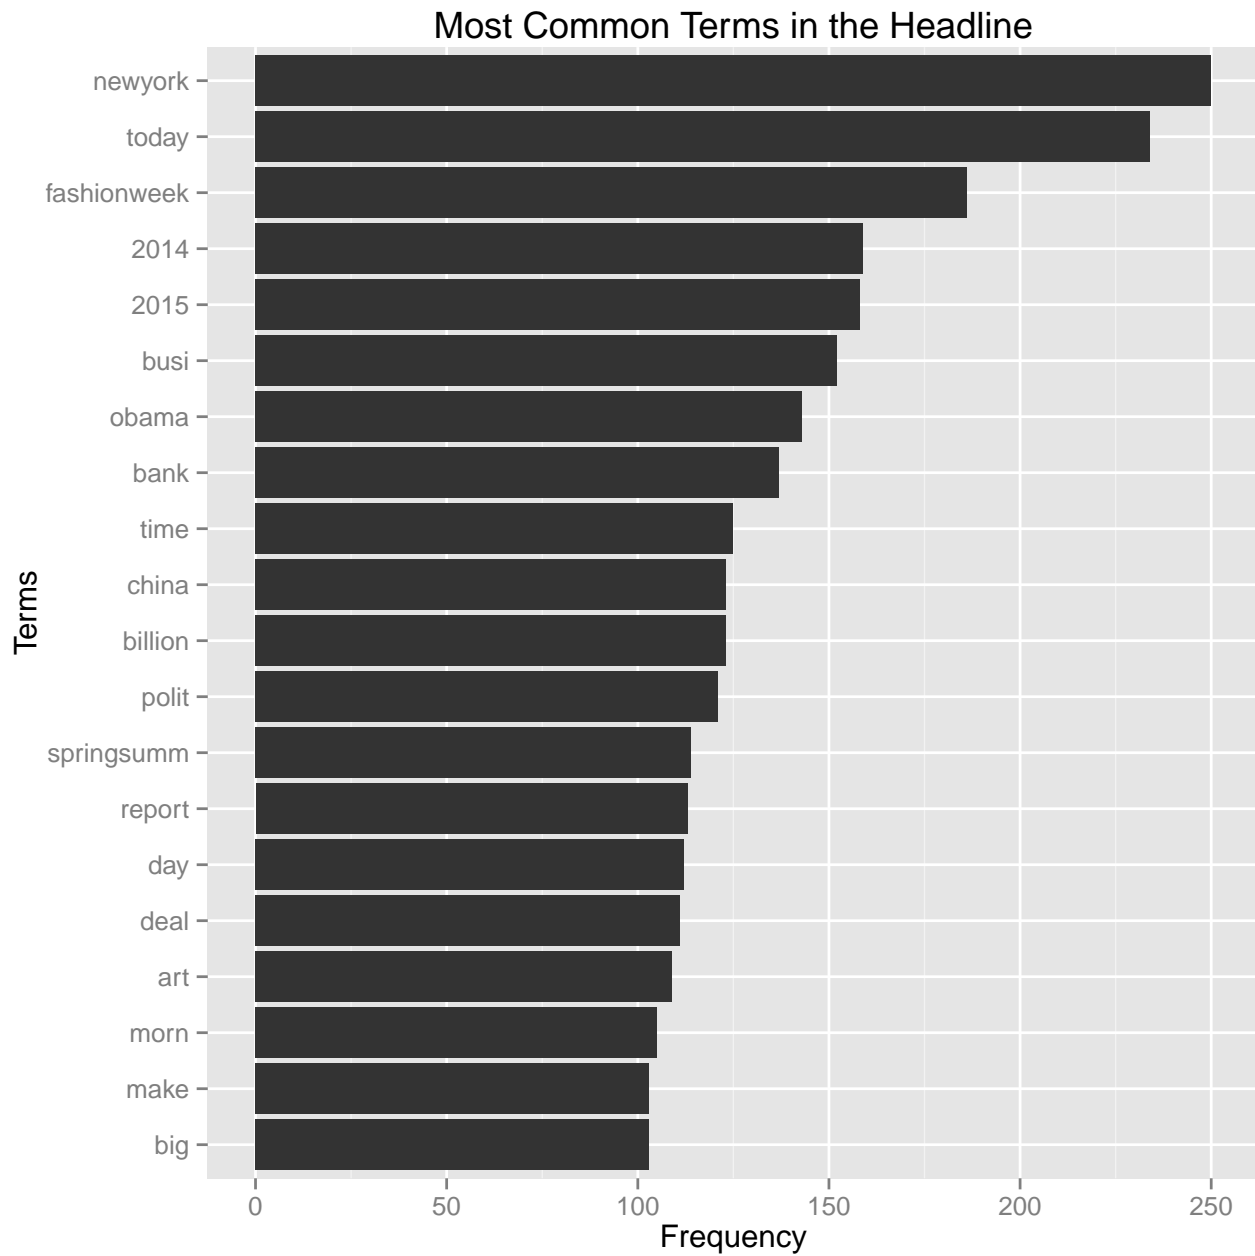
tdmText = TermDocumentMatrix(CorpusText)
# around 20 words
sparseText = removeSparseTerms(tdmText, 0.989)
sparseText = as.data.frame(as.matrix(sparseText))
colnames(sparseText) = make.names(colnames(sparseText))

dtmText = DocumentTermMatrix(CorpusText)
freqTerms = findFreqTerms(dtmText, lowfreq=10)
termFreq = colSums(as.matrix(dtmText))
termFreq = subset(termFreq, termFreq>=100)
df = data.frame(term=names(termFreq), freq=termFreq)

## Bag of Words
newsDataBoW = newsData
newsDataBoW$PubDate = NULL

tSparseText = t(sparseText)
colnames(tSparseText) = make.names(paste('c', colnames(tSparseText), sep="_"))
newsDataBoW[, colnames(tSparseText)] = tSparseText

```

Modeling Data

Logistic Regression without contents feature

```
# MODEL--0
# LR_0
# finding importance of features
removedColumns = c("SectionName", "NewsDesk", "SubsectionName", "Headline", "Snippet", "Abstract", "Summary")
LR_0 = glm(Popular ~ ., data=newsTrain[,!colnames(newsTrain) %in% removedColumns], family=binomial)

calcAUCLr(LR_0, newsTrain$Popular)
```

```
## [1] 0.9015615 0.9315814
```

```

# training result: 0.9315814
LR_0_Pred = predict(LR_0, newdata=newsTest[, !colnames(newsTrain) %in% removedColumns], type="response")

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

generateSubmission(LR_0_Pred, submitPath)
# testing result: 0.89088

summary(LR_0)

```

```

##
## Call:
## glm(formula = Popular ~ ., family = binomial, data = newsTrain[,
##      !colnames(newsTrain) %in% removedColumns])
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.8005  -0.3645  -0.1895  -0.0059   3.4498
##
## Coefficients: (12 not defined because of singularities)
##
##              Estimate Std. Error z value
## (Intercept)      3.885e+00  4.792e+02   0.008
## PubDay           1.131e-04  1.940e-03   0.058
## Hour             3.572e-02  9.917e-03   3.601
## WordCount        3.052e-04  1.462e-04   2.088
## HeadlineCharCount -1.409e-02  8.095e-03  -1.741
## SummaryCharCount  -1.544e-02  4.066e-03  -3.797
## HeadlineWordCount  1.080e-01  4.170e-02   2.589
## SummaryWordCount   7.890e-02  2.350e-02   3.357
## LogWordCount      8.855e-01  1.073e-01   8.253
## ShortHeadline1     8.897e-02  1.647e-01   0.540
## DayOfWeek         -5.135e-01  3.607e-01  -1.424
## DayOfWeekTuesday   1.740e-01  4.397e-01   0.396
## DayOfWeekWednesday 8.961e-01  7.822e-01   1.146
## DayOfWeekThursday  1.272e+00  1.128e+00   1.128
## DayOfWeekFriday    1.625e+00  1.500e+00   1.083
## DayOfWeekSaturday  2.324e+00  2.123e+00   1.095
## DayOfWeekSunday    NA          NA          NA
## Holiday1          -9.462e-02  2.385e-01  -0.397
## SectionNameFactorBusiness Day -1.397e+01  4.781e+02  -0.029
## SectionNameFactorHealth -1.547e+01  6.523e+03  -0.002
## SectionNameFactorMagazine -1.593e+00  3.993e-01  -3.990
## SectionNameFactorMissing  2.868e-01  1.159e+00   0.247
## SectionNameFactorMultimedia -2.263e+00  1.379e+00  -1.641
## SectionNameFactorN.Y. / Region 6.207e-02  3.405e-01   0.182
## SectionNameFactorOpen -1.582e+01  3.104e+03  -0.005
## SectionNameFactorOpinion  4.031e+00  2.187e-01  18.430
## SectionNameFactorPuzzles  3.368e+00  3.537e-01   9.523
## SectionNameFactorScience  2.554e+00  7.987e+03   0.000
## SectionNameFactorSports -1.827e+01  4.607e+03  -0.004
## SectionNameFactorStyle  1.583e+00  2.384e-01   6.641

```

## SectionNameFactorTechnology	7.135e-01	2.506e-01	2.847
## SectionNameFactorTravel	-1.809e+00	1.030e+00	-1.757
## SectionNameFactorU.S.	-1.520e+01	3.599e+03	-0.004
## SectionNameFactorWorld	-1.536e+01	4.781e+02	-0.032
## NewsDeskFactorCulture	NA	NA	NA
## NewsDeskFactorForeign	NA	NA	NA
## NewsDeskFactorMagazine	-1.407e+01	1.088e+03	-0.013
## NewsDeskFactorMetro	NA	NA	NA
## NewsDeskFactorMissing	5.706e-01	1.156e+00	0.494
## NewsDeskFactorNational	NA	NA	NA
## NewsDeskFactorOpEd	NA	NA	NA
## NewsDeskFactorScience	1.831e+01	6.523e+03	0.003
## NewsDeskFactorSports	NA	NA	NA
## NewsDeskFactorStyles	NA	NA	NA
## NewsDeskFactorTechnology	NA	NA	NA
## NewsDeskFactorTravel	NA	NA	NA
## NewsDeskFactorTStyle	NA	NA	NA
## SubsectionNameFactorDealbook	1.011e+00	4.785e-01	2.114
## SubsectionNameFactorEducation	-1.351e+01	3.646e+03	-0.004
## SubsectionNameFactorFashion & Style	-3.108e+01	4.636e+03	-0.007
## SubsectionNameFactorMissing	-1.330e+01	4.781e+02	-0.028
## SubsectionNameFactorPolitics	1.799e+00	3.631e+03	0.000
## SubsectionNameFactorRoom For Debate	-2.134e+01	4.781e+02	-0.045
## SubsectionNameFactorSmall Business	NA	NA	NA
## SubsectionNameFactorThe Public Editor	-1.369e+01	4.781e+02	-0.029
## NumDailyArticles	4.342e-03	4.175e-03	1.040
## NumDailySectionArticles	-4.238e-02	1.550e-02	-2.735
## NumHourlyArticles	-5.318e-02	1.820e-02	-2.922
##	Pr(> z)		
## (Intercept)	0.993531		
## PubDay	0.953512		
## Hour	0.000316 ***		
## WordCount	0.036821 *		
## HeadlineCharCount	0.081762 .		
## SummaryCharCount	0.000147 ***		
## HeadlineWordCount	0.009638 **		
## SummaryWordCount	0.000787 ***		
## LogWordCount	< 2e-16 ***		
## ShortHeadline1	0.589008		
## DayOfWeek	0.154528		
## DayOfWeekTuesday	0.692276		
## DayOfWeekWednesday	0.251948		
## DayOfWeekThursday	0.259461		
## DayOfWeekFriday	0.278668		
## DayOfWeekSaturday	0.273643		
## DayOfWeekSunday	NA		
## Holiday1	0.691587		
## SectionNameFactorBusiness Day	0.976684		
## SectionNameFactorHealth	0.998108		
## SectionNameFactorMagazine	6.60e-05 ***		
## SectionNameFactorMissing	0.804539		
## SectionNameFactorMultimedia	0.100866		
## SectionNameFactorN.Y. / Region	0.855372		
## SectionNameFactorOpen	0.995935		

```

## SectionNameFactorOpinion          < 2e-16 ***
## SectionNameFactorPuzzles          < 2e-16 ***
## SectionNameFactorScience          0.999745
## SectionNameFactorSports            0.996837
## SectionNameFactorStyle             3.12e-11 ***
## SectionNameFactorTechnology        0.004415 **
## SectionNameFactorTravel            0.078882 .
## SectionNameFactorU.S.             0.996631
## SectionNameFactorWorld             0.974370
## NewsDeskFactorCulture              NA
## NewsDeskFactorForeign              NA
## NewsDeskFactorMagazine             0.989684
## NewsDeskFactorMetro                NA
## NewsDeskFactorMissing              0.621492
## NewsDeskFactorNational             NA
## NewsDeskFactorOpEd                 NA
## NewsDeskFactorScience              0.997760
## NewsDeskFactorSports               NA
## NewsDeskFactorStyles               NA
## NewsDeskFactorTechnology           NA
## NewsDeskFactorTravel               NA
## NewsDeskFactorTStyle               NA
## SubsectionNameFactorDealbook       0.034544 *
## SubsectionNameFactorEducation      0.997043
## SubsectionNameFactorFashion & Style 0.994651
## SubsectionNameFactorMissing        0.977800
## SubsectionNameFactorPolitics       0.999605
## SubsectionNameFactorRoom For Debate 0.964404
## SubsectionNameFactorSmall Business NA
## SubsectionNameFactorThe Public Editor 0.977166
## NumDailyArticles                  0.298318
## NumDailySectionArticles            0.006239 **
## NumHourlyArticles                  0.003480 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 5900.1  on 6531  degrees of freedom
## Residual deviance: 3093.0  on 6486  degrees of freedom
## AIC: 3185
##
## Number of Fisher Scoring iterations: 17

```

By using `summary(Model)`, we can select the following variables based on importance of variables.

```

selectedColumns = c("Popular",
                    "PubDay",
                    "Hour",
                    "LogWordCount",
                    "SummaryCharCount",
                    "HeadlineWordCount",
                    "WordCount",
                    "SummaryCharCount",

```

```

"HeadlineCharCount",
"SectionNameFactor",
"SubsectionNameFactor",
"NumDailySectionArticles",
"NumHourlyArticles")

```

MODEL-1 (logistic regression)

```

LR_1 = glm(Popular ~ ., data=newsTrain[, colnames(newsTrain) %in% selectedColumns], family=binomial)

calcAUCLr(LR_1, newsTrain$Popular)
# training result: 0.9297705

LR_1_Pred = predict(LR_1, newdata=newsTest[, colnames(newsTrain) %in% selectedColumns], type="response")

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

generateSubmission(LR_1_Pred, submitPath)
# testing result: 0.89212

```

MODEL-2 (random forest)

```

RF_2 = randomForest(Popular ~ ., data=newsTrain[, colnames(newsTrain) %in% selectedColumns], nodesize=1

trainPartition = createDataPartition(y=newsTrain$Popular, p=0.5, list=FALSE)
tuneTrain      = newsTrain[trainPartition, colnames(newsTrain) %in% selectedColumns]
RF_2.tuned    = train(Popular ~ ., data=tuneTrain, method="rf", trControl=trainControl(method="cv", number=10))

calcAUC(RF_2, newsTrain$Popular)
# training result: 0.9387521

RF_2_Pred = predict(RF_2, newdata=newsTest, type="prob")[,2]
generateSubmission(RF_2_Pred, submitPath)
# testing result: 0.89470

```

MODEL-3 (logistic regression on Bag of Words)

```

LR_3 = glm(Popular ~ ., data=newsTrainBoW[, !colnames(newsTrainBoW) %in% removedColumns], family=binomial)

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

calcAUCLr(LR_3, newsTrainBoW$Popular)
# training result: 0.9389116

LR_3_Pred = predict(LR_3, newdata=newsTestBoW, type="response")

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

```

```
generateSubmission(LR_3_Pred, submitPath)
# testing result: 0.90176
```

MODEL-4 (random forest on Bag of Words)

```
RF_4 = randomForest(Popular ~ ., data=newsTrainBoW[,!colnames(newsTrainBoW) %in% removedColumns], nodes=
trainPartition = createDataPartition(y=newsTrainBoW$Popular, p=0.5, list=FALSE)
tuneTrain      = newsTrainBoW[trainPartition,!colnames(newsTrainBoW) %in% removedColumns]
RF_4.tuned     = train(Popular ~ ., data=tuneTrain, method="rf", trControl=trainControl(method="cv", number=
calCAUC(RF_4, newsTrainBoW$Popular)
# training result: 0.9399500


RF_4_Pred = predict(RF_4, newdata=newsTestBoW, type="prob")[,2]
generateSubmission(RF_4_Pred, submitPath)
# testing result: 0.90048
```

Supervised + Unsupervised

```
matrixSparseText = as.matrix(sparseText)
matrixSparseText.distMatrix = dist(scale(matrixSparseText))
matrixSparseText.clusters = hclust(matrixSparseText.distMatrix, method="ward.D2")

dText = as.dendrogram(matrixSparseText.clusters)
dTextData <- dendro_data(dText, type = "rectangle")

ggplot(segment(dTextData)) +
  geom_segment(aes(x = x, y = y, xend = xend, yend = yend)) + coord_flip()
```



```
kText = 25
mText = t(sparseText)
KMCText = kmeans(mText, kText)

for (i in 1:kText) {
  cat(paste("cluster", i, ": ", sep=" "))
  s = sort(KMCText$centers[i, ], decreasing=TRUE)
  cat(names(s)[1:15], sep=", ", "\n")
}

newsData$TextCluster = as.factor(KMCText$cluster)
newsData$PubDate = NULL
newsTrain = head(newsData, nrow(trainData))
newsTest = tail(newsData, nrow(testData))
```

MODEL-5 (random forest on text cluster)

```
selectedColumns = cbind(selectedColumns, c("TextCluster"))
RF_5 = randomForest(Popular ~ ., data=newsTrain[, colnames(newsTrain) %in% selectedColumns], nodesize=5

trainPartition = createDataPartition(y=newsTrain$Popular, p=0.5, list=FALSE)
tuneTrain = newsTrain[trainPartition, colnames(newsTrain) %in% selectedColumns]
RF_5.tuned = train(Popular ~ ., data=tuneTrain, method="rf", trControl=trainControl(method="cv", number=5))

calcAUC(RF_5, newsTrain$Popular)
# training result: 0.9399861
```

```
RF_5_Pred = predict(RF_5, newdata=newsTest, type="prob")[,2]
generateSubmission(RF_5_Pred, submitPath)
# testing result: 0.89854
```

Features of negative example

```
newsData$QAboutNews = as.factor(ifelse(grepl("6q", newsData$Text) == TRUE, 1, 0))
newsData$DailyClip = as.factor(ifelse(grepl("DailyClipReport", newsData$Text) == TRUE, 1, 0))
newsData$DailyReport = as.factor(ifelse(grepl("DailyReport", newsData$Text) == TRUE, 1, 0))
newsData$PicOfDay = as.factor(ifelse(grepl("PicOfDay", newsData$Text) == TRUE, 1, 0))
newsData$TestYourself = as.factor(ifelse(grepl("TestYourself", newsData$Text) == TRUE, 1, 0))
newsData$WordOfDay = as.factor(ifelse(grepl("WordOfDay", newsData$Text) == TRUE, 1, 0))
newsData$Recap = as.factor(ifelse(grepl("recap", newsData$Text) == TRUE, 1, 0))

newsTrain = head(newsData, nrow(trainData))
newsTest = tail(newsData, nrow(testData))
```

MODEL-6 (logistic regression with negative example features)

```
selectedColumns = cbind(selectedColumns, c("QAboutNews", "DailyClip", "DailyReport", "PicOfDay", "TestYourself", "WordOfDay", "Recap"))

LR_6 = glm(Popular ~ ., data=newsTrain[, colnames(newsTrain) %in% selectedColumns], family=binomial)
calcAUClr(LR_6, newsTrain$Popular)

LR_6_Pred = predict(LR_6, newdata=newsTest[, colnames(newsTest) %in% selectedColumns], type="response")
generateSubmission(LR_6_Pred, submitPath)
```

MODEL-7 (random forest with negative example features)

```
RF_7 = randomForest(Popular ~ ., data=newsTrain[, colnames(newsTrain) %in% selectedColumns], nodesize=5)

trainPartition = createDataPartition(y=newsTrain$Popular, p=0.5, list=FALSE)
tuneTrain = newsTrain[trainPartition, colnames(newsTrain) %in% selectedColumns]

RF_7.tuned = train(Popular ~ ., data=tuneTrain, method="rf", trControl=trainControl(method="cv", number=10))

calcAUC(RF_7, newsTrain$Popular)
# training result: 0.9390981

RF_7_Pred = predict(RF_7, newdata=newsTest, type="prob")[,2]
generateSubmission(RF_7_Pred, submitPath)
# testing result: 0.89867
```

Topic Words

Add some topic features according [Topic from Google](#).

```
questionWords <- c("\\?", "^why", "should", "^when", "can", "^if", "^is")
religionWords <- c("secular", "humanist", "humanism", "secularist", "god", "religion", "atheist", "atheism", "islamists", "church", "atheists", "jesus", "christ", "christian", "catholic", "pope", "imam", "\\")
techWords <- c("apple", "\\<ios\\>", "ipod", "ipad", "iphone")
```



```

healthWords <- c("cancer", "weight", "fat", "heart", "disease", "brain", "sex", "sexual", "love", "hate", "doctor")
sciWords <- c("climate", "warming", "global", "science", "scientists")
busWords <- c("jobs", "employment", "work", "working", "economy")
poliHiWords <- c("republican", "conservative")

newsData$Question <- as.factor(ifelse(grepl(paste(questionWords, collapse="|"), newsData$Headline)==TRUE, 1, 0))
newsData$religionWords <- ifelse(grepl(paste(religionWords, collapse="|"), newsData$Headline)==TRUE, 1, 0)
newsData$tech <- ifelse(grepl(paste(techWords, collapse="|"), newsData$Headline)==TRUE, 1, 0)
newsData$health <- ifelse(grepl(paste(healthWords, collapse="|"), newsData$Headline)==TRUE, 1, 0)
newsData$sci <- ifelse(grepl(paste(sciWords, collapse="|"), newsData$Headline)==TRUE, 1, 0)
newsData$business <- ifelse(grepl(paste(busWords, collapse="|"), newsData$Headline)==TRUE, 1, 0)
newsData$poliHi <- ifelse(grepl(paste(poliHiWords, collapse="|"), newsData$Headline)==TRUE, 1, 0)
newsData$noComment <- ifelse(grepl("no comment necessary", newsData$Headline), 1, 0)
newsData$comments <- ifelse(grepl("open for comments", newsData$Headline), 1, 0)

```

MODEL-8 (logistic regression with topic words)

```
LR_8 = glm(Popular ~ . - UniqueID, data=newsTrain, family=binomial)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```

calcAUClr(LR_8, newsTrain$Popular)
# training result: 0.9412723

```

```
LR_8_Pred = predict(LR_8, newdata=newsTest, type="response")
```

```

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

```

```

generateSubmission(LR_8_Pred, submitPath)
# testing result: 0.89517

```

MODEL-9 (random forest with topic words)

```

RF_9 = randomForest(Popular ~ . - UniqueID, data=newsTrain, nodesize=5, ntree=2500, importance=TRUE)
trainPartition = createDataPartition(y=newsTrain$Popular, p=0.5, list=FALSE)
tuneTrain      = newsTrain[trainPartition, ]
RF_9.tuned    = train(Popular ~ . - UniqueID, data=tuneTrain, method="rf", trControl=trainControl(method=

calcAUC(RF_9, newsTrain$Popular)
# training result: 0.9418562

RF_9_Pred = predict(RF_9, newdata=newsTest, type="prob")[,2]
generateSubmission(RF_9_Pred, submitPath)
# testing result: 0.89962

```

Model Evaluation

Model Name	Train ROC	Test ROC	DESC
LR_0	0.9315814	0.89088	logistic regression

Model Name	Train ROC	Test ROC	DESC
LR_1	0.9297705	0.89212	logistic regression on selected variables
LR_3	0.9389116	0.90176	logistic regression with bag of words
LR_6	0.9367199	0.90108	logistic regression with negative features
LR_8	0.9411255	0.89801	logistic regression with topic words
RF_2	0.9388249	0.89368	random forest
RF_4	0.9399030	0.89965	random forest with bag of words
RF_5	0.9387642	0.89865	random forest on text cluster
RF_7	0.9389404	0.89732	random forest with negative features
RF_9	0.941084	0.89887	random forest with topic words

Ensemble

```

EN_1 = (2*RF_9_Pred + 3*RF_4_Pred + 1*LR_3_Pred)/6
EN_2 = (2*RF_4_Pred + 3*LR_3_Pred + 1*LR_3_Pred)/6
EN_3 = (2*LR_1_Pred + 3*RF_2_Pred + 1*LR_8_Pred)/6
EN_4 = (2*LR_6_Pred + 3*RF_5_Pred + 1*RF_7_Pred)/6
EN_5 = (2*RF_5_Pred + 3*RF_4_Pred + 1*LR_3_Pred)/6
EN_6 = (1*RF_4_Pred + 2*LR_3_Pred + 1*LR_3_Pred)/4

generateSubmission(EN_1, submitPath)
generateSubmission(EN_2, submitPath)
generateSubmission(EN_3, submitPath)
generateSubmission(EN_4, submitPath)
generateSubmission(EN_5, submitPath)
generateSubmission(EN_6, submitPath)

```

Model Name	Test ROC
EN_1	0.90238
EN_2	0.90550
EN_3	0.89973
EN_4	0.90361
EN_5	0.90271
EN_6	0.90529